

Metric definitions

Latency — per-request `latency_ms`, rolling median and p95.

Target drift — share of recommendations per *sub-genre* over time. A recommender has no predicted classes; sub-genre (from the catalog category path) is the stated proxy.

Live accuracy — fraction of user feedback marked 👍 (a precision proxy), reported with sample size n.


 Refresh data




Wayfinder Ops — model monitoring

Total predictions	Median latency	Fallback rate	Live accuracy 
22	0 ms	59%	54%

Prediction latency over time



Recommended sub-genre distribution (target-drift proxy)



Share of recommended books per sub-genre, hourly. A stable mix = no drift; a shifting mix = the model's output distribution is moving.

Live accuracy (user feedback)



Cumulative share of 👍 across 26 feedback events. Interpret with n in view — early values rest on few votes.

model_version	accuracy	n
philosophy-recommender:v0	0.5385	26